Transformers

1. A Transformer Network, like its predecessors RNNs, GRUs and LSTMs, can process information one word at a time. (Sequential architecture). **1 / 1 point**

- (•) False
- ( ) True

Expand

Correct
Correct! A Transformer Network can ingest entire sentences all at the same time.

2. The major innovation of the transformer architecture is combining the use of LSTMs and RNN sequential processing. **1 / 1 point**

- ( ) True
- (•) False

Expand

Correct
The major innovation of the transformer architecture is combining the use of attention based representations and a CNN convolutional neural network style of processing.

3. The concept of *Self-Attention* is that: **1 / 1 point**

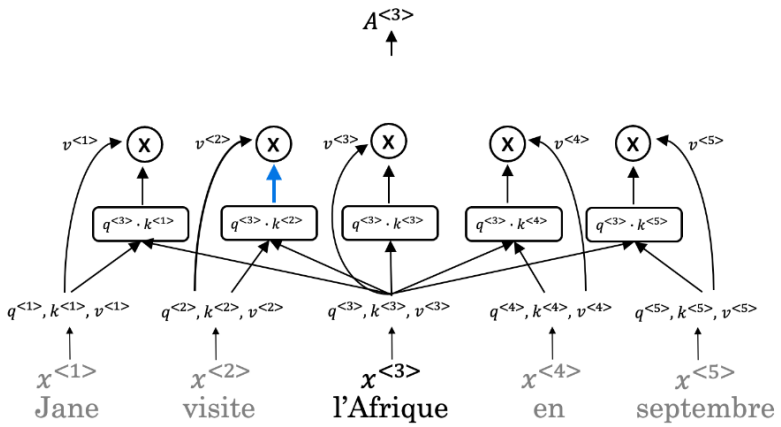

- Given a word, its neighbouring words are used to compute its context by selecting the highest of those word values to map the Attention related to that given word.
- Given a word, its neighbouring words are used to compute its context by taking the average of those word values to map the Attention related to that given word.
- Given a word, its neighbouring words are used to compute its context by summing up the word values to map the Attention related to that given word.
- Given a word, its neighbouring words are used to compute its context by selecting the lowest of those word values to map the Attention related to that given word.

Expand

Correct

4. What letter does the "?" represent in the following representation of *Attention*? **1 / 1 point**

$$Attention(Q, K, V) = softmax(\frac{QK^T}{\sqrt{d_?}})V$$

- (•) k
- ( ) q
- ( ) v
- ( ) t

Expand

Correct
k is represented by the $?$ in the representation.

5. Which of the following statements represents Key (K) as used in the self-attention calculation? **1 / 1 point**

- ( ) K = the order of the words in a sentence
- ( ) K = specific representations of words given a Q
- (•) K = qualities of words given a Q
- ( ) K = interesting questions about the words in a sentence

Expand

Correct
The qualities of words given a Q are represented by Key (K).

6. $Attention(W_i^Q Q, W_i^K K, W_i^V V)$ 1 / 1 point

What does i represent in this multi-head attention computation?

- The computed attention weight matrix associated with specific representations of words given a Q
- The computed attention weight matrix associated with the *ith* "head" (sequence)
- The computed attention weight matrix associated with the *ith* "word" in a sentence.
- The computed attention weight matrix associated with the order of the words in a sentence

Expand

Correct

i here represents the computed attention weight matrix associated with the "head" (sequence).

7. Following is the architecture within a Transformer Network ***(without displaying positional encoding and output layers(s)).***

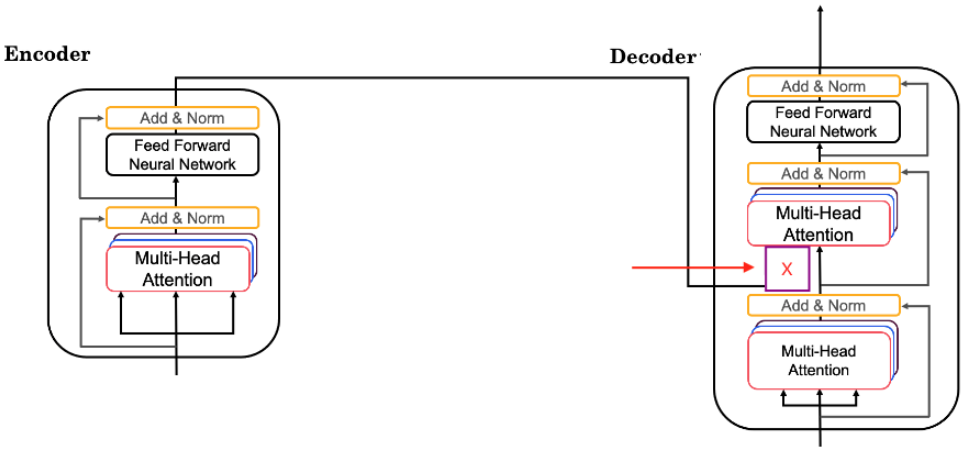

What information does the *Decoder* take from the *Encoder* for its second block of *Multi-Head Attention*? (Marked X, pointed by the independent arrow)

What information does the *Decoder* take from the *Encoder* for its second block of *Multi-Head Attention*? (Marked X, pointed by the independent arrow)

(Check all that apply)

- [ ] Q
- [x] V

Correct

- [x] K

Correct

Expand

Correct
Great, you got all the right answers.

8. Following is the architecture within a Transformer Network. ***(without displaying positional encoding and output layers(s))*** **1 / 1 point**

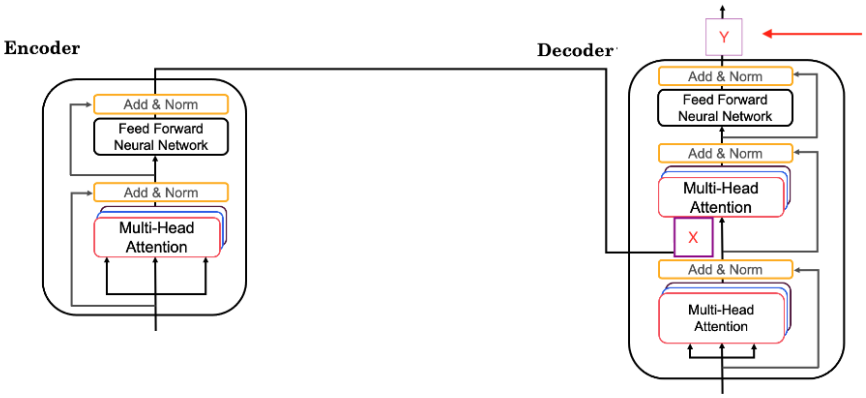

What is the output layer(s) of the *Decoder*? (Marked Y, pointed by the independent arrow)

What is the output layer(s) of the *Decoder*? (Marked Y, pointed by the independent arrow)

- (●) Linear layer followed by a softmax layer.
- ( ) Linear layer
- ( ) Softmax layer followed by a linear layer.
- ( ) Softmax layer

Expand

Correct

9. Which of the following statements is true about positional encoding? Select all that apply. **0 / 1 point**

- [ ] Positional encoding is used in the transformer network and the attention model.
- [x] Positional encoding is important because position and word order are essential in sentence construction of any language.

Correct
This is a correct answer, but other options are also correct. To review the concept watch the lecture *Transformer Network*.

- [ ] Positional encoding uses a combination of sine and cosine equations.
- [x] Positional encoding provides extra information to our model.

Correct
This is a correct answer, but other options are also correct. To review the concept watch the lecture *Transformer Network*.

Expand

Incorrect
You didn't select all the correct answers

10. Which of these is a good criterion for a good positional encoding algorithm? **1 / 1 point**

- ○ It must be nondeterministic.
- ○ It should output a common encoding for each time-step (word's position in a sentence).
- ◉ The algorithm should be able to generalize to longer sentences.
- ○ Distance between any two time-steps should be inconsistent for all sentence lengths.

Expand

Correct
This is a good criterion for a good positional encoding algorithm.

1. A Transformer Network, like its predecessors RNNs, GRUs and LSTMs, can process information one word at a time. (Sequential architecture). **1 / 1 point**

- ◉ False
- ○ True

Expand

Correct
Correct! A Transformer Network can ingest entire sentences all at the same time.

2. Transformer Network methodology is taken from: **1 / 1 point**

- (●) Attention Mechanism and CNN style of processing.
- ( ) Attention Mechanism and RNN style of processing.
- ( ) RNN and LSTMs
- ( ) GRUs and LSTMs

Expand

Correct
Transformer architecture combines the use of attention based representations and a CNN convolutional neural network style of processing.

3. The concept of *Self-Attention* is that: **1 / 1 point**

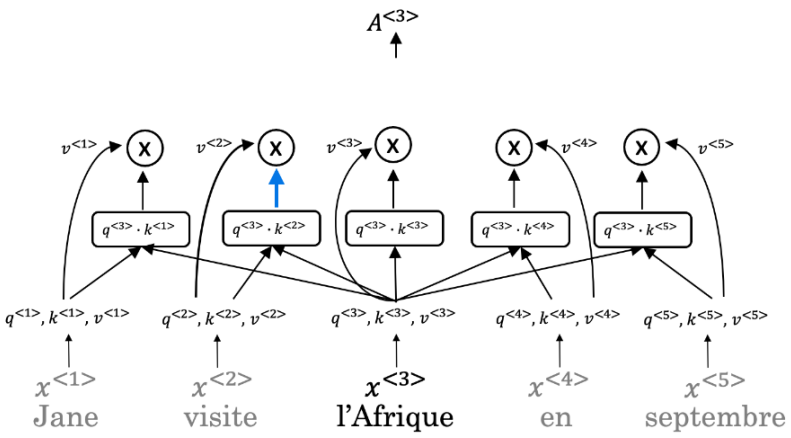

- Given a word, its neighbouring words are used to compute its context by selecting the highest of those word values to map the Attention related to that given word.
- Given a word, its neighbouring words are used to compute its context by summing up the word values to map the Attention related to that given word.
- Given a word, its neighbouring words are used to compute its context by taking the average of those word values to map the Attention related to that given word.
- Given a word, its neighbouring words are used to compute its context by selecting the lowest of those word values to map the Attention related to that given word.

Expand

Correct

4. Which of the following correctly represents *Attention ?*

1 / 1 point

- $Attention(Q, K, V) = softmax(\frac{QK^T}{\sqrt{d_k}})V$
- $Attention(Q, K, V) = min(\frac{QK^T}{\sqrt{d_k}})V$
- $Attention(Q, K, V) = min(\frac{QV^T}{\sqrt{d_k}})K$

Expand

Correct

5. Are the following statements true regarding Query (Q), Key (K) and Value (V)? 1 / 1 point

Q = interesting questions about the words in a sentence

K = qualities of words given a Q

V = specific representations of words given a Q

- ○ False
- ◉ True

Expand

Correct
Q = interesting questions about the words in a sentence, K = qualities of words given a Q, V = specific representations of words given a Q

6. $Attention(W_i^Q Q, W_i^K K, W_i^V V)$ 1 / 1 point

What does i represent in this multi-head attention computation?

- ○ The computed attention weight matrix associated with the order of the words in a sentence
- ○ The computed attention weight matrix associated with specific representations of words given a Q
- ○ The computed attention weight matrix associated with the *ith* "word" in a sentence.
- ◉ The computed attention weight matrix associated with the *ith* "head" (sequence)

Expand

Correct
i here represents the computed attention weight matrix associated with the "head" (sequence).

7. Following is the architecture within a Transformer Network ***(without displaying positional encoding and output layers(s)).*** 0 / 1 point

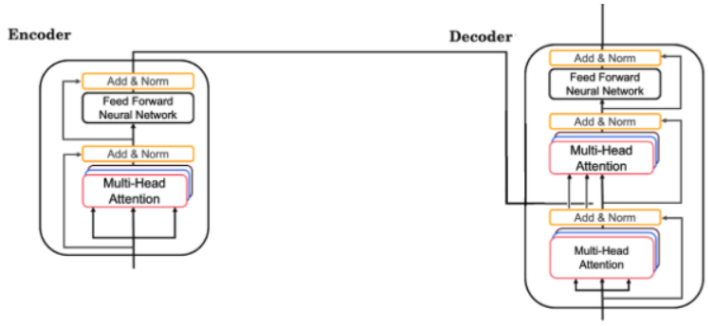

What is generated from the output of the *Decoder's* first block of *Multi-Head Attention*?

- (●) V
- ( ) Q
- ( ) K

Expand

Incorrect

To revise the concept watch the lecture .

8. Following is the architecture within a Transformer Network ***(without displaying positional encoding and output layers(s)).*** 1 / 1 point

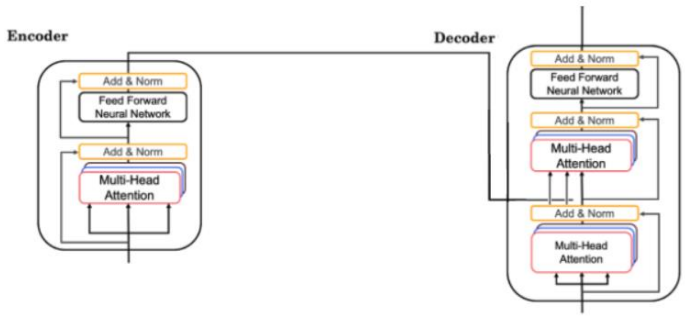

The output of the decoder block contains a softmax layer followed by a linear layer to predict the next word one word at a time.

The output of the decoder block contains a softmax layer followed by a linear layer to predict the next word one word at a time.

- ○ True
- ◉ False

Expand

Correct

The output of the decoder block contains a linear layer followed by a softmax layer to predict the next word one word at a time.

9. Which of the following statements is true? **0 / 1 point**

- ○ The transformer network differs from the attention model in that only the transformer network contains positional encoding.
- ○ The transformer network is similar to the attention model in that neither contain positional encoding.
- ○ The transformer network differs from the attention model in that only the attention model contains positional encoding.
- ◉ The transformer network is similar to the attention model in that both contain positional encoding.

Expand

⊗ **Incorrect**
To revise the concept watch the lecture .

10. Which of these is ***not*** a good criterion for a good positional encoding algorithm? **0 / 1 point**

- ○ It should output a common encoding for each time-step (word's position in a sentence).
- ○ Distance between any two time-steps should be consistent for all sentence lengths.
- ○ The algorithm should be able to generalize to longer sentences.
- ◉ It must be deterministic.

Expand

⊗ **Incorrect**
This is a good criterion for a good positional encoding algorithm.